

Application Layer

Jim KUROSE
www-aml.cs.umass.edu

Derived from Chapter 2 of

J. F. Kurose, K. W. Ross, "Computer Networking: A Top-Down Approach
Featuring the Internet", Addison Wesley Longman, 2001
<http://ocwonline.pearsoned.com/bookbind/pubbooks/kurose-ross1/chapter0/deluxe.html>

Chapter 2: Application Layer

Chapter goals:

- conceptual + implementation aspects of network application protocols
 - client server paradigm
 - service models
- learn about protocols by examining popular application-level protocols

More chapter goals

- specific protocols:
 - http
 - ftp
 - smtp
 - pop
 - dns
- programming network applications
 - socket programming


Applications and application-layer protocols

Application: communicating, distributed processes

- running in network hosts in "user space"
- exchange messages to implement app
- e.g., email, file transfer, the Web

Application-layer protocols

- one "piece" of an app
- define messages exchanged by apps and actions taken
- user services provided by lower layer protocols



Client-server paradigm


Typical network app has two pieces: *client* and *server*

Client:

- initiates contact with server ("speaks first")
- typically requests service from server,
- e.g.: request WWW page, send email

Server:

- provides requested service to client
- e.g., sends requested WWW page, receives/stores received email



Application-layer protocols (cont).

API: application programming interface

- defines interface between application and transport layer
- socket: Internet API
 - two processes communicate by sending data into socket, reading data out of socket

Q: how does a process "identify" the other process with which it wants to communicate?

- IP address of host running other process
- "port number" - allows receiving host to determine to which local process the message should be delivered

... lots more on this later.

What transport service does an app need?

Data loss

- some apps (e.g., audio) can tolerate some loss
- other apps (e.g., file transfer, telnet) require 100% reliable data transfer

Bandwidth

- some apps (e.g., multimedia) require minimum amount of bandwidth to be "effective"
- other apps ("elastic apps") make use of whatever bandwidth they get

Timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be "effective"

Transport service requirements of common apps

Application	Data loss	Bandwidth	Time Sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	loss-tolerant	elastic	no
real-time audio/video	loss-tolerant	audio: 5Kb-1Mb video: 10Kb-5Mb	yes, 100's msec
stored audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	few Kbps up	yes, 100's msec
financial apps	no loss	elastic	yes and no

Services provided by Internet transport protocols

TCP service:

- *connection-oriented*: setup required between client, server
- *reliable transport* between sending and receiving process
- *flow control*: sender won't overwhelm receiver
- *congestion control*: throttle sender when network overloaded
- *does not providing*: timing, minimum bandwidth guarantees

UDP service:

- unreliable data transfer between sending and receiving process
- does not provide: connection setup, reliability, flow control, congestion control, timing, or bandwidth guarantee

Q: why bother? Why is there a UDP?


Internet apps: their protocols and transport protocols

Application	Application layer protocol	Underlying transport protocol
e-mail	smtp [RFC 821]	TCP
remote terminal access	telnet [RFC 854]	TCP
Web	http [RFC 2068]	TCP
file transfer	ftp [RFC 959]	TCP
streaming multimedia	proprietary (e.g. RealNetworks)	TCP or UDP
remote file server	NSF	TCP or UDP
Internet telephony	proprietary (e.g., Vocaltec)	typically UDP

WWW: the http protocol

http: hypertext transfer protocol

- WWW's application layer protocol
- client/server model
 - *client*: browser that requests, receives, "displays" WWW objects
 - *server*: WWW server sends objects in response to requests
- http1.0: RFC 1945
- http1.1: RFC 2068



The http protocol: more

http: TCP transport service:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- http messages (application-layer protocol messages) exchanged between browser (http client) and WWW server (http server)
- TCP connection closed

http is "stateless"

- server maintains no information about past client requests

aside

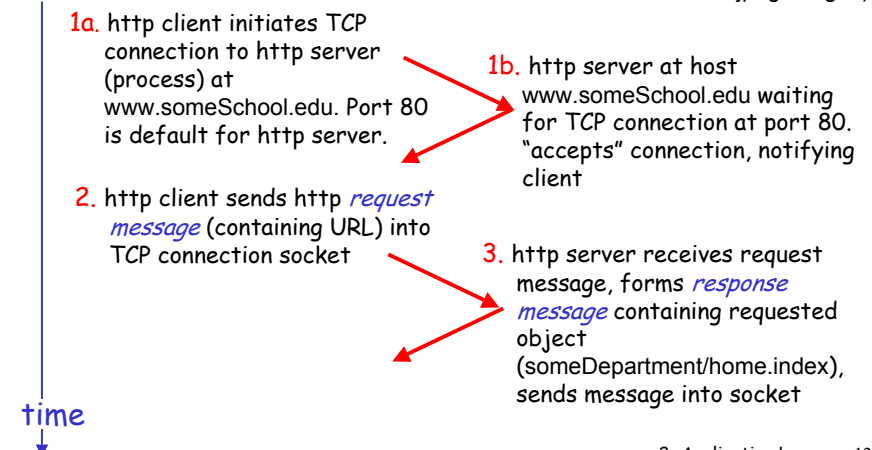
Protocols that maintain "state" are complex!

- past history (state) must be maintained
- if server/client crashes, their views of "state" may be inconsistent, must be reconciled

http example

Suppose user enters URL

www.someSchool.edu/someDepartment/home.index (contains text, references to 10 jpeg images)



http example (cont.)

- time ↓
4. http server closes TCP connection.
 5. http client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects
 6. Steps 1-5 repeated for each of 10 jpeg objects
- **non-persistent connection:** one object in each TCP connection
 - some browsers create multiple TCP connections *simultaneously* - one per object
 - **persistent connection:** multiple objects transferred within one TCP connection

http message format: request

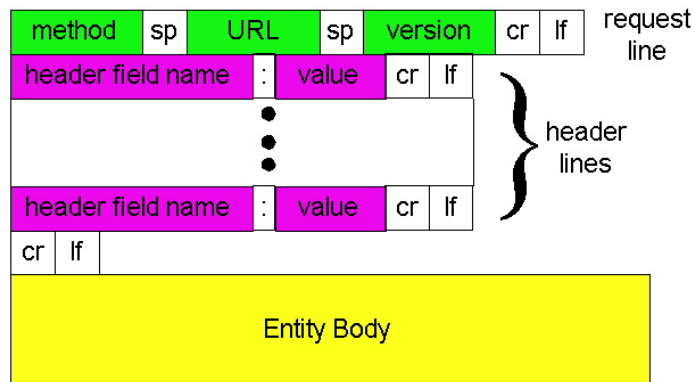
- two types of http messages: *request, response*
- **http request message:**
 - ASCII (human-readable format)

request line (GET, POST, HEAD commands) → GET /somedir/page.html HTTP/1.1

header lines → Connection: close
User-agent: Mozilla/4.0
Accept: text/html, image/gif, image/jpeg
Accept-language: fr

Carriage return line feed (extra carriage return, line feed) indicates end of message

http request message: general format



http message format: reply

status line (protocol status code status phrase) → HTTP/1.1 200 OK

header lines → Connection: close
Date: Thu, 06 Aug 1998 12:00:15 GMT
Server: Apache/1.3.0 (Unix)
Last-Modified: Mon, 22 Jun 1998
Content-Length: 6821
Content-Type: text/html

data, e.g., requested html file → data data data data data ...

http reply status codes

In first line in server->client response message.

A few sample codes:

200 OK

- request succeeded, requested object later in this message

301 Moved Permanently

- requested object moved, new location specified later in this message (Location:)

400 Bad Request

- request message not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

Trying out http (client side) for yourself

1. Telnet to your favorite WWW server:

```
telnet www.eurecom.fr 80
```

Opens TCP connection to port 80 (default http server port) at www.eurecom.fr. Anything typed in sent to port 80 at www.eurecom.fr

2. Type in a GET http request:

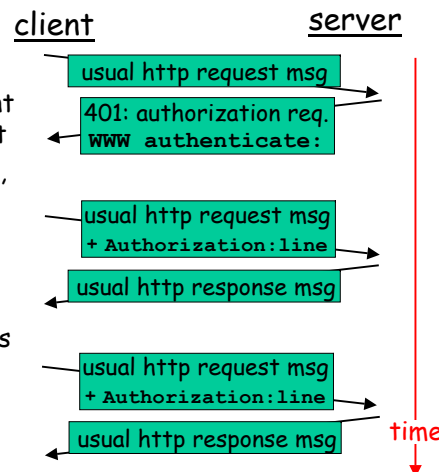
```
GET /~ross/index.html HTTP/1.0
```

By typing this in (hit carriage return twice), you send this minimal (but complete) GET request to http server

3. Look at response message sent by http server!

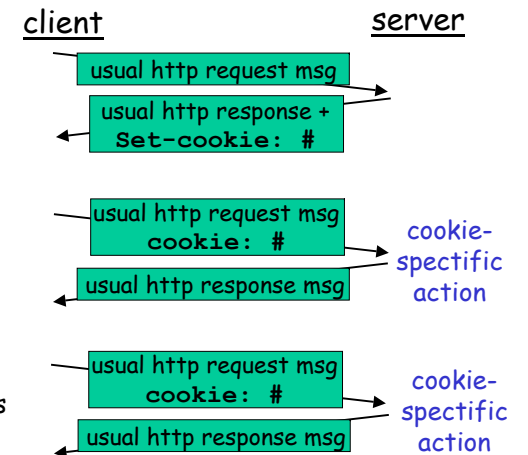
User-server interaction: authentication

- Authentication goal:** control access to server documents
- **stateless:** client must present authorization in each request
 - **authorization:** typically name, password
 - **authorization:** header line in request
 - if no authorization presented, server refuses access, sends
`WWW authenticate:` header line in response

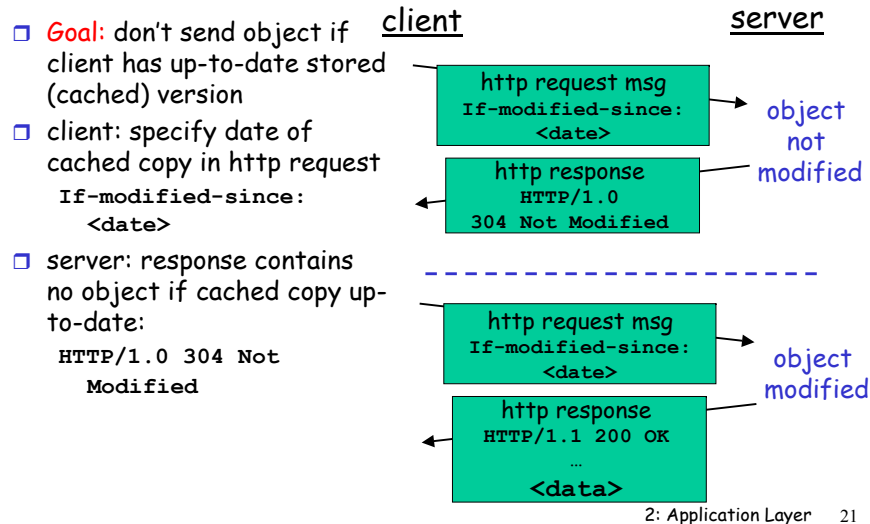


User-server interaction: cookies

- server sends "cookie" to client in response
`Set-cookie: #`
- client present cookie in later requests
`cookie: #`
- server matches presented-cookie with server-stored cookies
 - authentication
 - remembering user preferences, previous choices




User-server interaction: conditional GET



Web Caches (proxy server)

Goal: satisfy client request without involving origin server

- user sets browser: WWW accesses via web cache
- client sends all http requests to web cache
 - if object at web cache, web cache immediately returns object in http response
 - else requests object from origin server, then returns http response to client




2: Application Layer 22

Why WWW Caching?


Assume: cache is "close" to client (e.g., in same network)

- smaller response time: cache "closer" to client
- decrease traffic to distant servers
 - link out of institutional/local ISP network often bottleneck



2: Application Layer 23

ftp: the file transfer protocol




- transfer file to/from remote host
- client/server model
 - client:** side that initiates transfer (either to/from remote)
 - server:** remote host
- ftp: RFC 959
- ftp server: port 21

2: Application Layer 24

ftp: separate control, data connections

- ftp client contacts ftp server at port 21, specifying TCP as transport protocol
- two parallel TCP connections opened:
 - control:** exchange commands, responses between client, server. "out of band control"
 - data:** file data to/from server
- ftp server maintains "state": current directory, earlier authentication



ftp commands, responses

Sample commands:

- sent as ASCII text over control channel
- USER** *username*
- PASS** *password*
- LIST** return list of file in current directory
- RETR** *filename* retrieves (gets) file
- STOR** *filename* stores (puts) file onto remote host

Sample return codes

- status code and phrase (as in http)
- 331 Username OK, password required
- 125 data connection already open; transfer starting
- 425 Can't open data connection
- 452 Error writing file


Electronic Mail

Three major components:

- user agents
- mail servers
- simple mail transfer protocol: smtp

User Agent


- a.k.a. "mail reader"
- composing, editing, reading mail messages
- e.g., Eudora, pine, elm, Netscape Messenger
- outgoing, incoming messages stored on server



Electronic Mail: mail servers

Mail Servers

- mailbox** contains incoming messages (yet ot be read) for user
- message** queue of outgoing (to be sent) mail messages
- smtp protocol** between mail server to send email messages
 - client: sending mail server
 - "server": receiving mail server



Electronic Mail: smtp [RFC 821]

- ❑ uses tcp to reliably transfer email msg from client to server, port 25
- ❑ direct transfer: sending server to receiving server
- ❑ three phases of transfer
 - handshaking (greeting)
 - transfer
 - closure
- ❑ command/response interaction
 - **commands**: ASCII text
 - **response**: status code and phrase

Sample smtp interaction

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C:   How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

smtp: final words

try smtp interaction for yourself:

- ❑ telnet servername 25
- ❑ see 220 reply from server
- ❑ enter HELO, MAIL FROM, RCPT TO, DATA, QUIT commands

above lets you send email without using email client (reader)

Comparison with http


- ❑ http: pull
- ❑ email: push
- ❑ both have ASCII command/response interaction, status codes
- ❑ http: multiple objects in file sent in separate connections
- ❑ smtp: multiple message parts sent in one connection

Mail message format

smtp: protocol for exchanging email msgs

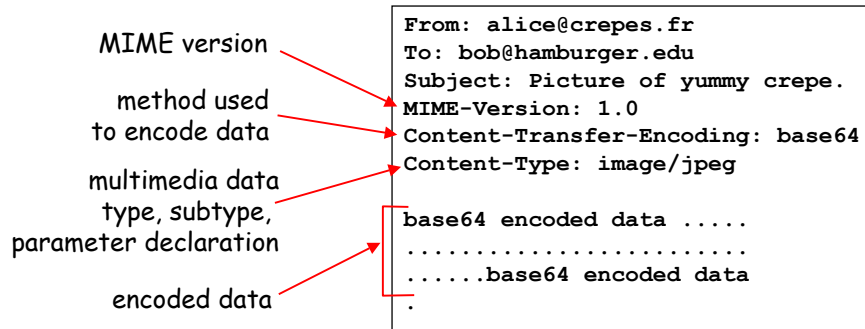
RFC 822: standard for text message format:

- ❑ header lines, e.g.,
 - To:
 - From:
 - Subject:*different from smtp commands!*
- ❑ body
 - the "message", ASCII characters only
- ❑ line containing only `.`



Message format: multimedia extensions

- MIME: multimedia mail extension, RFC 2045, 2056
- additional lines in msg header declare MIME content type



2: Application Layer 33

MIME types

Text

- example subtypes: plain, html

Video

- example subtypes: mpeg, quicktime

Image

- example subtypes: jpeg, gif

Application


- other data that must be processed by reader before "viewable"
- example subtypes: msword, octet-stream

Audio

- example subtypes: basic (8-bit mu-law encoded), 32kadpcm (32 kbps coding)

2: Application Layer 34

Mail access protocols



- SMTP: delivery/storage to receiver's server
- Mail access protocol: retrieval from server
 - POP: Post Office Protocol [RFC 1939]
 - authorization (agent <--> server) and download
 - IMAP: Internet Mail Access Protocol [RFC 1730]
 - more features (more complex)
 - manipulation of stored msgs on server

2: Application Layer 35

POP3 protocol

authorization phase

- client commands:
 - user: declare username
 - pass: password
- server responses
 - +OK
 - -ERR

transaction phase, client:

- list: list message numbers
- retr: retrieve message by number
- dele: delete
- quit

```
S: +OK POP3 server ready
C: user alice
S: +OK
C: pass hungry
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

2: Application Layer 36